

```
from google.colab import files
uploaded = files.upload()
```

Choose Files Dataset.csv
Dataset.csv(text/csv) - 1911016 bytes, last modified: 2/10/2026 - 100% done
 Saving Dataset.csv to Dataset.csv

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv('Dataset.csv', encoding='utf-8')
df
```

	Text	Language	
0	Nature, in the broadest sense, is the natural...	English	
1	"Nature" can refer to the phenomena of the phy...	English	
2	The study of nature is a large, if not the onl...	English	
3	Although humans are part of nature, human acti...	English	
4	[1] The word nature is borrowed from the Old F...	English	
...	
10332	ನಿಮ್ಮ ತಪ್ಪು ಏನು ಬಂದಿದೆಯೆಂದರೆ ಆ ದಿನದಿಂದ ನಿಮಗೆ ಒ...	Kannada	
10333	ನಾಸೀನ್‌ನಾ ತಾನು ಮೌರಲೀಗೆ ಹೆಚ್ಚಾಡುತ್ತಿದ್ದ ಮಾರ್ಗಗಳನ್ನ...	Kannada	
10334	ಹೇಗೆ 'ನಾಸೀನ್‌ನಿಸವ್' ಈಗ ಮರಿಯನ್ ಅವರಿಗೆ ಸಂಭವಿಸಿದ ಎ...	Kannada	
10335	ಅವಳು ಈಗ ಹೆಚ್ಚು ಚಿನ್ನದ ಬ್ರೆಕ್ ಬಯಸುವುದಿಲ್ಲ, ಎಂದು ...	Kannada	
10336	ಟೆರ್ಫ್ ನೀವು ನಿಜವಾಗಿಯೂ ಆ ದೇವದೂತನೆಂತೆ ಸ್ವಲ್ಪ ಕಾಣಿ...	Kannada	

10337 rows × 2 columns

Next steps: [Generate code with df](#) [New interactive sheet](#)

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, LSTM, Embedding, Dense, Bidirectional, Dropout
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.model_selection import train_test_split
```

```
df = pd.read_csv("Dataset.csv")

texts = df['Text'].astype(str).values
labels = df['Language'].astype('category').cat.codes.values
```

```
df['Language'].unique()
```

```
array(['English', 'Malayalam', 'Hindi', 'Tamil', 'Portugeese', 'French',
       'Dutch', 'Spanish', 'Greek', 'Russian', 'Danish', 'Italian',
       'Turkish', 'Swedish', 'Arabic', 'German', 'Kannada'], dtype=object)
```

```
tokenizer = Tokenizer(char_level=True)
tokenizer.fit_on_texts(texts)
seq = tokenizer.texts_to_sequences(texts)
padded_seq = pad_sequences(seq, maxlen=150)
```

```
X_train, X_test, y_train, y_test = train_test_split(
    padded_seq, labels, test_size=0.2, random_state=42
)
```

```
input_layer = Input(shape=(150,)) # increased length

embed = Embedding(
    input_dim=len(tokenizer.word_index)+1,
    output_dim=128
)(input_layer)

x = Bidirectional(LSTM(128, return_sequences=True))(embed)
x = Dropout(0.3)(x)

x = Bidirectional(LSTM(64))(x)
x = Dropout(0.3)(x)

x = Dense(64, activation='relu')(x)

output = Dense(len(np.unique(labels)), activation='softmax')(x)

model = Model(inputs=input_layer, outputs=output)
```

```
model.compile(
    optimizer='adam',
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

```
from tensorflow.keras.callbacks import EarlyStopping

early_stop = EarlyStopping(
    monitor='val_loss',
    patience=3,
    restore_best_weights=True
)
```

```
history = model.fit(
    X_train, y_train,
    epochs=20,
    batch_size=32,
    validation_split=0.1,
    callbacks=[early_stop]
)
```

```
Epoch 1/20
233/233 184s 756ms/step - accuracy: 0.2440 - loss: 2.1936 - val_accuracy: 0.5091 -
Epoch 2/20
233/233 176s 757ms/step - accuracy: 0.5671 - loss: 1.1124 - val_accuracy: 0.5877 -
Epoch 3/20
233/233 168s 722ms/step - accuracy: 0.6647 - loss: 0.8615 - val_accuracy: 0.6542 -
Epoch 4/20
233/233 201s 716ms/step - accuracy: 0.7326 - loss: 0.6811 - val_accuracy: 0.7570 -
Epoch 5/20
233/233 209s 744ms/step - accuracy: 0.7914 - loss: 0.5617 - val_accuracy: 0.7400 -
Epoch 6/20
233/233 172s 740ms/step - accuracy: 0.8090 - loss: 0.5094 - val_accuracy: 0.8235 -
Epoch 7/20
233/233 177s 759ms/step - accuracy: 0.8479 - loss: 0.4220 - val_accuracy: 0.8537 -
Epoch 8/20
233/233 172s 737ms/step - accuracy: 0.8776 - loss: 0.3570 - val_accuracy: 0.8863 -
```

```

Epoch 9/20
233/233 172s 740ms/step - accuracy: 0.8914 - loss: 0.3071 - val_accuracy: 0.9081 -
Epoch 10/20
233/233 172s 739ms/step - accuracy: 0.9109 - loss: 0.2550 - val_accuracy: 0.9178 -
Epoch 11/20
233/233 171s 734ms/step - accuracy: 0.9306 - loss: 0.2111 - val_accuracy: 0.9238 -
Epoch 12/20
233/233 170s 729ms/step - accuracy: 0.9308 - loss: 0.2086 - val_accuracy: 0.9299 -
Epoch 13/20
233/233 170s 727ms/step - accuracy: 0.9462 - loss: 0.1584 - val_accuracy: 0.9093 -
Epoch 14/20
233/233 175s 754ms/step - accuracy: 0.9362 - loss: 0.1921 - val_accuracy: 0.9420 -
Epoch 15/20
233/233 170s 730ms/step - accuracy: 0.9495 - loss: 0.1471 - val_accuracy: 0.9383 -
Epoch 16/20
233/233 170s 731ms/step - accuracy: 0.9535 - loss: 0.1376 - val_accuracy: 0.9335 -
Epoch 17/20
233/233 175s 753ms/step - accuracy: 0.9580 - loss: 0.1212 - val_accuracy: 0.9432 -
Epoch 18/20
233/233 175s 752ms/step - accuracy: 0.9617 - loss: 0.1128 - val_accuracy: 0.9553 -
Epoch 19/20
233/233 170s 732ms/step - accuracy: 0.9664 - loss: 0.0905 - val_accuracy: 0.9504 -
Epoch 20/20
233/233 170s 730ms/step - accuracy: 0.9720 - loss: 0.0762 - val_accuracy: 0.9432 -

```

```
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.1)
```

```

Epoch 1/10
233/233 170s 731ms/step - accuracy: 0.9645 - loss: 0.0995 - val_accuracy: 0.9528 -
Epoch 2/10
233/233 170s 728ms/step - accuracy: 0.9713 - loss: 0.0835 - val_accuracy: 0.9347 -
Epoch 3/10
233/233 169s 727ms/step - accuracy: 0.9602 - loss: 0.1215 - val_accuracy: 0.9625 -
Epoch 4/10
233/233 171s 732ms/step - accuracy: 0.9686 - loss: 0.0860 - val_accuracy: 0.9613 -
Epoch 5/10
233/233 203s 738ms/step - accuracy: 0.9732 - loss: 0.0790 - val_accuracy: 0.9553 -
Epoch 6/10
233/233 172s 737ms/step - accuracy: 0.9796 - loss: 0.0567 - val_accuracy: 0.9637 -
Epoch 7/10
233/233 171s 733ms/step - accuracy: 0.9808 - loss: 0.0529 - val_accuracy: 0.9613 -
Epoch 8/10
233/233 173s 739ms/step - accuracy: 0.9867 - loss: 0.0365 - val_accuracy: 0.9649 -
Epoch 9/10
233/233 171s 735ms/step - accuracy: 0.9878 - loss: 0.0348 - val_accuracy: 0.9637 -
Epoch 10/10
233/233 172s 740ms/step - accuracy: 0.9846 - loss: 0.0505 - val_accuracy: 0.9480 -
<keras.src.callbacks.history.History at 0x7832a846f500>

```

```

def predict_language(text):
    seq = tokenizer.texts_to_sequences([text])
    pad = pad_sequences(seq, maxlen=150)
    pred = model.predict(pad)
    lang_index = np.argmax(pred)
    return df['Language'].astype('category').cat.categories[lang_index]

print(predict_language("Hola amigo"))

```

```
1/1 1s 676ms/step
Spanish
```

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

y_pred_probs = model.predict(X_test)
y_pred = np.argmax(y_pred_probs, axis=1)

# Accuracy
acc = accuracy_score(y_test, y_pred)

```

```
# Precision, Recall, F1 (macro because multi-class)
prec = precision_score(y_test, y_pred, average='macro')
rec = recall_score(y_test, y_pred, average='macro')
f1 = f1_score(y_test, y_pred, average='macro')

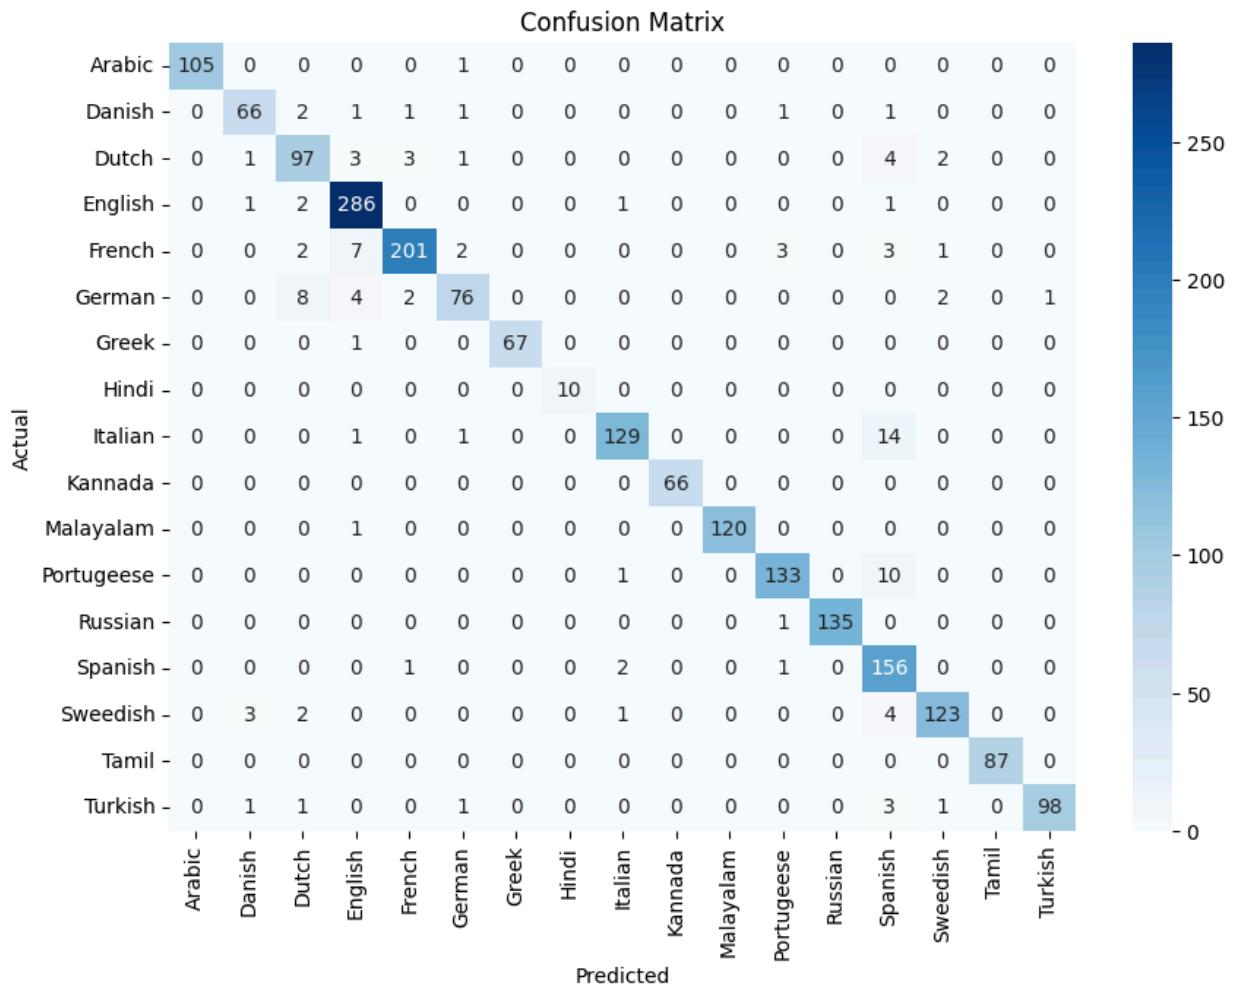
print("Accuracy :", acc)
print("Precision:", prec)
print("Recall   :", rec)
print("F1 Score :", f1)
```

65/65 _____ 13s 196ms/step
 Accuracy : 0.945357833655706
 Precision: 0.9558334686270075
 Recall : 0.9472040147596514
 F1 Score : 0.9505761412196032

```
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(10,7))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=df['Language'].astype('category').cat.categories,
            yticklabels=df['Language'].astype('category').cat.categories)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



```
model.save("language_identification_model.h5")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(mo

```
import pickle

with open("tokenizer.pkl", "wb") as f:
    pickle.dump(tokenizer, f)
```

```
labels_list = list(df['Language'].astype('category').cat.categories)

with open("labels.pkl", "wb") as f:
    pickle.dump(labels_list, f)
```

```
from tensorflow.keras.models import load_model
import pickle

model = load_model("language_identification_model.h5")

with open("tokenizer.pkl", "rb") as f:
    tokenizer = pickle.load(f)

with open("labels.pkl", "rb") as f:
    labels_list = pickle.load(f)
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_m

```
from google.colab import files

files.download("language_identification_model.h5")
files.download("tokenizer.pkl")
files.download("labels.pkl")
```

```
print(predict_language("नमस्ते दोस्त"))

print(predict_language("எப்படி இருக்கிறீர்கள்"))
print(predict_language("Hello World"))
print(predict_language("olá mundo"))
print(predict_language("Привет, мир"))
```

```
1/1 ━━━━━━━━━━ 1s 631ms/step
Malayalam
1/1 ━━━━━━ 0s 82ms/step
Tamil
1/1 ━━━━ 0s 82ms/step
English
1/1 ━━ 0s 84ms/step
Spanish
1/1 ━ 0s 85ms/step
Russian
```

```
def predict_lang(new_text):
    seq = tokenizer.texts_to_sequences([new_text])
    pad = pad_sequences(seq, maxlen=100)
    pred = model.predict(pad)
    idx = np.argmax(pred)
    return labels_list[idx]
```

```
print(predict_lang("Bonjour"))
```

1/1 ————— 0s 69ms/step
Spanish

```
import matplotlib.pyplot as plt

# Accuracy Graph
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'])
plt.show()

# Loss Graph
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'])
plt.show()
```

